

BP6: Enterprise AML Compliance Monitoring & Regulatory Reporting

Platform-Wide Compliance-Impact Rollup -- generated 2026-10-02T14:54:50.340394+00:00 UTC

Business Objective

BP6 (Enterprise AML Compliance Monitoring & Regulatory Reporting) is this platform's Phase 4 capstone: a pure, read-only executive rollup of BP1-BP5's own already-computed real figures, mirroring the oversight function JPMorgan Chase's own official Global Financial Crimes Compliance page names as its "AML Compliance Program" (per the locked master-execution-plan Section 3). BP6 trains no model and introduces no independent ground-truth label -- its entire real output traces back to BP1-BP5's own real MODEL_CARD.md/RULE_CARD.md Financial Impact/Rule Details sections and their own real Notebook-3 validation_report JSON files, bridged through one real, deterministic extractor script and this BP's own 4-notebook phase-then-platform aggregation pipeline.

Recommendation & Status

RECOMMENDED FOR PRODUCTION -- PLATFORM-WIDE

Every real business problem's own two-gate validation (BP1, BP2, BP3, BP4, BP5) passed on its own locked mandatory realism-validation tier. This platform-wide status is a pure pass-through of those five real, already-recorded verdicts -- nothing is re-validated or re-computed at the rollup level.

Executive Summary

This platform-wide rollup is a pure, read-only aggregation of BP1-BP5's own already-computed real figures. Sourcing method (verbatim): "Every figure below was parsed verbatim from a real, already-generated BP1-BP5 report file (MODEL_CARD.md / RULE_CARD.md Financial Impact / Rule Details section, and each BP's own real NB3 validation_report JSON). No figure here is recomputed, estimated, or fabricated -- this script performs pure extraction of already-computed real output, per the locked master plan's Section 8 'recomputing nothing' rule and Section 7A's 'Pure sum of BP1-BP5's own already-computed deltas' rule for BP6."

What This Means for the Business

- Platform-wide status: RECOMMENDED FOR PRODUCTION -- PLATFORM-WIDE. Every real business problem's own two-gate validation (BP1, BP2, BP3, BP4, BP5) passed on its own locked mandatory realism-validation tier. This platform-wide status is a pure pass-through of those five real, already-recorded verdicts -- nothing is re-validated or re-computed at the rollup level.
- Real BENEFIT category 1 (FP-reduction savings, BP1+BP4+BP5, ASSUMPTION-labeled): \$13,877,192 (\$13.88M) across 213,496 investigator hours and 853,981 fewer false-positive alerts.
- Real BENEFIT category 2 (TP-uplift illustrative value, BP1+BP4+BP5, ASSUMPTION-labeled): \$97,950,000 (\$97.95M) across +1,959 additional real cases caught -- never summed with category 1 above.
- Real BENEFIT category 3 (BP2's own two typology-detection figures, kept separate): \$1,313 auto-typing efficiency savings and \$1,104,000 (\$1.10M) typology-confirmation value.
- BP4 contributes the single largest real category-2 dollar figure of the platform's three contributing BPs (\$65,000,000 (\$65.00M)) -- and 'XGBoost' recurs as the real champion on 3 BPs (BP1, BP4, BP5).
- BP3's real network coverage (2,032,095 nodes, 4,363,197 edges) is reported as a portfolio-scale volume line -- never dollarized or summed with the BENEFIT categories above, consistent with this platform's locked never-blend-categories discipline.

Three-Category Financial Rollup (never blended)

This platform reports THREE separate BENEFIT lines (categories 1, 2, and 3 above) -- never one blended grand total -- per the locked Section 7A/Section 8 never-summed-across-categories discipline. A reader wanting a single number must choose which category answers their question; none is privileged as 'the' headline by this report.

Category	Line Item	Value	Note
BENEFIT 1 -- FP-Reduction Savings (ASSUMPTION; never summed with categories 2/3)	BP1 -- Transaction Monitoring	\$4,792,239 (\$4.79M)	FP-reduction \$ savings
BENEFIT 1 -- FP-Reduction Savings (ASSUMPTION; never summed with categories 2/3)	BP4 -- Structuring & Smurfing	\$4,232,784 (\$4.23M)	FP-reduction \$ savings
BENEFIT 1 -- FP-Reduction Savings (ASSUMPTION; never summed with categories 2/3)	BP5 -- Correspondent Banking & Cross-Border	\$4,852,169 (\$4.85M)	FP-reduction \$ savings

categories 2/3)			
BENEFIT 1 -- FP-Reduction Savings (ASSUMPTION; never summed with categories 2/3)	TOTAL (BP1+BP4+BP5)	\$13,877,192 (\$13.88M)	213,496 investigator hours saved, 853,981 fewer FP alerts
BENEFIT 2 -- TP-Uplift Illustrative Value (ASSUMPTION; never summed with categories 1/3)	BP1 -- Transaction Monitoring	\$15,950,000 (\$15.95M)	TP-uplift illustrative \$
BENEFIT 2 -- TP-Uplift Illustrative Value (ASSUMPTION; never summed with categories 1/3)	BP4 -- Structuring & Smurfing	\$65,000,000 (\$65.00M)	TP-uplift illustrative \$
BENEFIT 2 -- TP-Uplift Illustrative Value (ASSUMPTION; never summed with categories 1/3)	BP5 -- Correspondent Banking & Cross-Border	\$17,000,000 (\$17.00M)	TP-uplift illustrative \$
BENEFIT 2 -- TP-Uplift Illustrative Value (ASSUMPTION; never summed with categories 1/3)	TOTAL (BP1+BP4+BP5)	\$97,950,000 (\$97.95M)	+1,959 additional real cases caught
BENEFIT 3 -- BP2 Typology Lines (kept separate -- different unit basis, never blended into 1/2)	Auto-typing efficiency savings	\$1,313	BP2's own figure -- different case subset/ASSUMPTION basis than categories 1/2
BENEFIT 3 -- BP2 Typology Lines (kept separate -- different unit basis, never blended into 1/2)	Typology confirmation value	\$1,104,000 (\$1.10M)	BP2's own figure -- different case subset/ASSUMPTION basis than categories 1/2
COST CONTEXT (informational only -- never summed into any headline figure)	Platform build/operating cost	NOT SOURCED -- none invented	Informational only, never summed into any headline figure (locked Section 8 rule). This platform has no real sourced build/operating-

			cost figure for any BP -- none is invented here; this section intentionally carries no dollar figure, consistent with this platform's zero-fabrication discipline.
PORTFOLIO SCALE (volume only -- never dollarized, never summed with BENEFIT)	BP3 real network nodes	2,032,095	Structurally distinct from BP1/2/4/5's transaction-level populations
PORTFOLIO SCALE (volume only -- never dollarized, never summed with BENEFIT)	BP3 real network edges	4,363,197	Structurally distinct from BP1/2/4/5's transaction-level populations

Per-BP Status Grid

BP	Name	Primary Variant	Champion / Signal	Verdict
BP1	Transaction Monitoring & Suspicious Activity Detection	LI-Medium	XGBoost	PASS
BP2	Typology & Red-Flag Pattern Detection	LI-Medium	RandomForest	PASS
BP3	Transaction Network & Graph Intelligence	LI-Medium	2-hop proximity to a TRAIN-flagged account	PASS
BP4	Structuring & Smurfing Detection	LI-Medium	XGBoost	PASS
BP5	Correspondent Banking & Cross-Border Wire Risk	LI-Medium	XGBoost	PASS

Phase-Grouped Rollup

Phase 1 -- Detection Foundation

BPs: BP1, BP2

- BP1 real champion XGBoost (verdict PASS): \$4,792,239 FP-reduction savings, \$15,950,000 TP-uplift illustrative value.
- BP2 real champion RandomForest (verdict PASS): \$1,313 auto-typing efficiency savings, \$1,104,000 typology-confirmation value (kept on its own lines -- different unit basis than BP1's figures above).

Phase 2 -- Network & Specialized Typology

BPs: BP3, BP4

- BP3 real champion signal '2-hop proximity to a TRAIN-flagged account' (verdict PASS): 2,032,095 real network nodes, 4,363,197 real edges -- volume only, never dollarized (no sourced per-account cost basis).
- BP4 real champion XGBoost (verdict PASS): \$4,232,784 FP-reduction savings, \$65,000,000 TP-uplift illustrative value.

Phase 3 -- Cross-Border Intelligence

BPs: BP5

- BP5 real champion XGBoost (verdict PASS): \$4,852,169 FP-reduction savings, \$17,000,000 TP-uplift illustrative value.

Per-BP Business Narratives

What each business problem actually does, its real validated performance, and what deploying it to production would concretely mean operationally -- every sentence below cites a real number already present in this platform's own source data.

BP1 -- Transaction Monitoring & Suspicious Activity Detection

Verdict: PASS

What it does

BP1 is the real function financial institutions themselves call a 'transaction monitoring system' -- language drawn directly from a real Federal Reserve consent order against American Express Bank International -- and the 'computer monitoring system [that] issued alerts' language from FinCEN's own civil money penalty assessment against JPMorgan Chase. It scores every real transaction against this dataset's real, direct ground-truth label ('Is Laundering') using supervised imbalanced binary classification, with PR-AUC as the primary metric because laundering transactions are a small real fraction of all activity. This evidences USA PATRIOT Act Section 326/314(a)/(b) obligations and feeds FinCEN SAR filing (31 CFR 1020.320) -- it does not itself predict SAR-filed/not-filed.

Real validated performance

On LI-Medium -- this BP's locked mandatory realism-validation tier -- the real champion (XGBoost) passed both validation gates (verdict PASS) using 19 real engineered features, at a real selected decision threshold of 0.9798. At that operating point it achieves real precision 0.142 and real recall 0.179 (real PR-AUC 0.1242).

What production deployment would mean

Deployed to production, BP1's real measured improvement over a naive fixed-dollar-threshold rule means investigators stop reviewing 294,907 real false-positive alerts (ASSUMPTION-estimated at 73,727 investigator hours, \$4,792,239 (\$4.79M)), while the same real operating point independently catches 319 more real laundering cases than the naive rule would (illustrative regulatory-exposure-avoidance ASSUMPTION of \$15,950,000 (\$15.95M)) -- freeing real investigator capacity for genuine cases instead of chasing noise, and strengthening the evidence base feeding the SAR-filing process.

Methodology & Regulatory Basis

BP1's own real Business Objective is to build the core transaction-monitoring capability every Bank Secrecy Act-regulated institution must operate under the FFIEC's five pillars -- scoring individual transactions for real money-laundering likelihood and feeding investigator alert review and any downstream SAR decision, while controlling the false-positive alert volume that consumes most AML teams' real capacity. The real champion (XGBoost) reaches this using 19 real engineered features at a real decision threshold of 0.9798. Real SHAP analysis shows the model is driven most by Payment Format and the sender's distinct-counterparty count to date. Regulatory basis: the Bank Secrecy Act (31 U.S.C. Section 5311) platform-wide, USA PATRIOT Act Section 326 (CIP) and Section 314(a)/(b) (BP1, BP5), and FinCEN's SAR filing requirement (31 CFR Section 1020.320).

BP2 -- Typology & Red-Flag Pattern Detection

Verdict: PASS

What it does

BP2 answers a different question than BP1: not just 'is this suspicious' but 'which of the real labeled laundering typologies does it match' -- a real multi-class target this dataset's own 'Patterns.txt' ground truth supports across up to 8 typologies (fan-out, fan-in, gather-scatter, scatter-gather, cycle, random, bipartite, stack). This is the 'red flags' language FinCEN's own AML guidance uses for typology indicators, and the specific terminology FinCEN's real civil penalty assessment against JPMorgan Chase used describing red-flag indicators the bank was found to have missed. BP2 evidences FATF's 40 Recommendations and feeds the same real SAR process BP1 does, with typology-level detail BP1's binary flag cannot provide.

Real validated performance

On LI-Medium, the real champion (RandomForest) passed both validation gates (verdict PASS) using 13 real engineered features, achieving a real macro-F1 of 0.444 across the

typologies with enough real matched examples this run (the two largest real classes: GATHER-SCATTER with 541 real examples).

What production deployment would mean

Deployed, BP2 would auto-classify 101 more real cases by typology than a single-typology heuristic rule can (ASSUMPTION-estimated at 20.2 auto-typing hours saved, \$1,313) -- freeing analysts from manual typology lookup -- while independently confirming the correct typology on 138 real cases the old single-typology heuristic rule would have missed outright (typology-confirmation value ASSUMPTION of \$1,104,000 (\$1.10M)). These two figures are kept deliberately separate from BP1/BP4/BP5's categories (different unit basis, different real case subset) rather than blended into one platform total.

Methodology & Regulatory Basis

BP2's own real Business Objective is to classify transactions already known to be laundering into the specific real money-laundering typology they exhibit -- directly matching the 'red flags' terminology used in real regulatory enforcement actions -- giving investigators a real, evidence-backed starting typology for case triage and SAR-narrative drafting rather than leaving every case to be typed manually. The real champion (RandomForest) reaches this using 13 real engineered features (multi-class argmax, no single decision threshold). Real SHAP analysis shows the model is driven most by the receiver's distinct-counterparty count to date and the sender's distinct-counterparty count to date. Regulatory basis: the Bank Secrecy Act (31 U.S.C. Section 5311) platform-wide, FinCEN's own 'red flags' typology guidance from its JPMorgan Chase assessment (Madoff/BLM case) -- the specific real-world terminology this BP's typology framing matches -- and the same SAR filing requirement (31 CFR Section 1020.320) BP1 feeds.

BP3 -- Transaction Network & Graph Intelligence

Verdict: PASS

What it does

BP3 has no trained ML model and needs none -- it asks a structural question BP1/BP2/BP4/BP5 cannot: is an account's real position in the transaction network itself informative, independent of its own transaction history? Unlike the other four BPs' real-world grounding, this specific graph-analytics framing was not independently confirmed as AMEX/JPMorgan-specific terminology in the primary documents checked during this platform's own feasibility review -- honestly flagged rather than overclaimed -- though network/graph analytics is standard industry AML practice, and this BP's real output carries GDPR cross-border-data-handling relevance alongside BP5.

Real validated performance

On LI-Medium -- a real graph of 2,032,095 nodes and 4,363,197 edges -- the real champion signal ('2-hop proximity to a TRAIN-flagged account') passed both validation gates (verdict PASS) with a real network-lift ratio of 3.19x over the real base rate of 1.157% (bootstrap 95% CI [3.157x, 3.219x], confirming the lift is real and not sampling noise).

What production deployment would mean

Deployed, BP3 would not replace BP1/BP2/BP4/BP5's own scoring -- it would hand investigators a real bounded 2-hop ego-network around every flagged account (illustrative real examples sized 181,831, 15,360, 15,174 real accounts before the locked 2,000-node cap) instead of reviewing that one account in isolation. There is deliberately no dollar figure here -- No sourced real per-account investigation-cost basis exists for this BP (disclosed honestly in the real RULE_CARD.md -- none is invented here). -- but a 3.19x real lift at 2,032,095-node platform scale is a materially different, network-aware triage signal than single-account review alone.

Methodology & Regulatory Basis

BP3's own real Business Objective is to flag accounts using directly-interpretable graph-structural signals -- real degree, real PageRank, and real network proximity to already-flagged accounts -- computed from the real account-to-account transaction structure, with no ground-truth label required to construct them. BP3 has no trained ML model and genuinely no SHAP/LIME explainability section (not merely omitted); the real champion signal ('2-hop proximity to a TRAIN-flagged account') is ranked purely by its real network-lift ratio on a real held-out test set of accounts. Regulatory basis: the Bank Secrecy Act (31 U.S.C. Section 5311) platform-wide and the FFIEC's five-pillar BSA/AML Examination Manual, under which link analysis and relationship mapping are a recognized investigative technique -- honestly flagged as not independently confirmed against a primary enforcement document, unlike the other four BPs' real-world grounding.

BP4 -- Structuring & Smurfing Detection

Verdict: PASS

What it does

BP4 targets a real named federal crime: 31 U.S.C. Section 5324 makes it illegal to break transactions into smaller pieces specifically to evade Bank Secrecy Act reporting requirements, and detecting this pattern is a mandatory examination area at every BSA-regulated institution. BP4 engineers 5 real, rule-derived structuring/smurfing features (sub-threshold amount clustering, time-clustering, fan-out splitting) on top of BP1's existing feature set, testing whether structuring-specific behavior carries real additional predictive lift for the same 'Is Laundering' label BP1 uses.

Real validated performance

On LI-Medium, the real champion (XGBoost) passed both validation gates (verdict PASS) using 24 real features (BP1's 19 plus BP4's own 5 structuring-specific ones) at a real threshold of 0.9793, achieving real precision 0.086 and real recall 0.305 (real PR-AUC 0.1253) -- the highest real recall of the platform's three binary BPs, at the cost of lower real precision, a real trade-off worth flagging for threshold review.

What production deployment would mean

Deployed, BP4 contributes the single largest real dollar figure on the entire platform: an illustrative regulatory-exposure-avoidance ASSUMPTION of \$65,000,000 (\$65.00M) from

1,300 additional real structuring cases caught that a naive rule would miss, plus \$4,232,784 (\$4.23M) in false-positive-reduction savings (65,120 investigator hours, 260,479 fewer false-positive alerts). Given structuring's own named statute and mandatory-examination status, this is real evidence this is the platform's highest-priority capability to keep in continuous production monitoring.

Methodology & Regulatory Basis

BP4's own real Business Objective targets a federal crime with its own named statute: 31 U.S.C. Section 5324 makes it illegal to break transactions into smaller pieces specifically to evade Bank Secrecy Act reporting requirements, and detecting this pattern is a mandatory examination area at every BSA-regulated institution. Beyond BP1's shared feature set, BP4 adds its own real 5 named structuring-specific features -- `structuring\_window\_txn\_count`, `structuring\_window\_amount\_sum`, `daily\_fanout\_count`, `amount\_to\_rolling\_window\_mean\_ratio`, `near\_threshold\_flag` -- of which `amount\_to\_rolling\_window\_mean\_ratio` ranks among the real top-5 global SHAP drivers this run: the model overall is driven most by Payment Format and the sender's distinct-counterparty count to date. Regulatory basis: BP4's own named statute, the BSA structuring statute (31 U.S.C. Section 5324), sits alongside the Bank Secrecy Act (31 U.S.C. Section 5311) platform-wide and the same SAR filing requirement (31 CFR Section 1020.320) BP1/BP2 feed.

BP5 -- Correspondent Banking & Cross-Border Wire Risk

Verdict: PASS

What it does

BP5 targets correspondent banking and cross-border wire risk -- the specific area that drove the real HSBC, Standard Chartered, and Danske Bank enforcement actions, and a standing high-priority area industry-wide. It uses this dataset's real cross-institution, cross-border structure (bank-to-bank transfer fields plus the real country-tagged bank names in the accounts data, e.g. 'Portugal Bank #4507') to score cross-border flow risk, evidencing USA PATRIOT Act Section 314(a)/(b), OFAC sanctions-list screening context, the Wolfsberg AML Principles for correspondent banking, and GDPR cross-border-data-handling relevance alongside BP3.

Real validated performance

On LI-Medium, the real champion (XGBoost) passed both validation gates (verdict PASS) using 27 real features -- the most of any binary BP on the platform, reflecting the added cross-border/country-pair signal -- at a real threshold of 0.9814, achieving real precision 0.176 and real recall 0.183 (real PR-AUC 0.1399).

What production deployment would mean

Deployed, BP5 delivers the platform's single largest real false-positive-reduction figure: \$4,852,169 (\$4.85M) in savings from 74,649 investigator hours freed across 298,595 fewer false-positive alerts -- more than either BP1 or BP4 -- plus \$17,000,000 (\$17.00M) illustrative regulatory-exposure-avoidance from 340 additional real cross-border cases

caught. Given the real enforcement history this capability is modeled against, keeping it in continuous production use directly addresses the exact risk area regulators have most recently and visibly penalized at other institutions.

Methodology & Regulatory Basis

BP5's own real Business Objective targets the single highest-risk channel for sanctioned-party exposure and cross-border money-laundering flow: wires crossing a national border. Beyond the platform's shared 19-feature transaction-monitoring set, BP5 adds its own real 8 named cross-border features -- `sender\_country`, `receiver\_country`, `is\_cross\_border`, `sender\_is\_foreign`, `receiver\_is\_foreign`, `foreign\_to\_foreign`, `sender\_country\_empirical\_risk`, `receiver\_country\_empirical\_risk`. None of these 8 cross-border-specific features broke into the real top-5 overall SHAP drivers this run -- the model overall remains driven most by Payment Format and the sender's distinct-counterparty count to date -- though this BP's own MODEL_CARD.md separately ranks `sender\_country\_empirical\_risk` as the top driver specifically among its own cross-border feature subset. Regulatory basis: USA PATRIOT Act Section 326 (Customer Identification Program) for correspondent-banking onboarding, Section 314(a)/(b) for cross-border information sharing -- this BP's own named obligation -- and the Wolfsberg Group's Correspondent Banking Due Diligence Questionnaire (CBDDQ) principles; OFAC sanctions-list screening is part of this BP's regulatory framing but is Not Possible - Data Limitation on this synthetic dataset.

Regulatory & Compliance Mapping

Framework	Applies to
Bank Secrecy Act (31 U.S.C. Section 5311)	Platform-wide
FFIEC BSA/AML Examination Manual (5 pillars)	Platform-wide
FATF 40 Recommendations	BP2, BP6
Wolfsberg AML Principles (correspondent banking)	BP5, BP6
SR 11-7 Model Risk Management	Every model-bearing BP (BP6 itself has no model -- pure rollup)

Fairness / Bias Testing

Not Possible - Data Limitation: the IBM AML transaction dataset carries no demographic fields, so an ECOA/FCRA-style disparate-impact test cannot be computed for any contributing BP, and therefore not for this platform-wide rollup either. Flagged explicitly here rather than silently skipped, per this platform's locked compliance-documentation policy.

Limitations & Caveats

- This platform-wide rollup is a pure pass-through of BP1-BP5's own already-computed real figures -- it introduces no new model, no new ground-truth label, and re-validates nothing beyond the real structural and reconciliation gates in this BP's own Notebook 3.
- Every dollar figure labeled ASSUMPTION traces back to that BP's own disclosed assumption set (see that BP's own Excel workbook Assumptions sheet) -- none is a specific real fine, penalty, or contractual figure.
- BENEFIT categories 1, 2, and 3, COST CONTEXT, and PORTFOLIO SCALE are never summed into one blended total anywhere in this package -- a reader wanting a single number must choose which category answers their question.
- No real sourced build/operating-cost figure exists for this platform -- COST CONTEXT carries no dollar figure, honestly, rather than an invented placeholder.
- Fairness/disparate-impact testing is not computable on the underlying IBM AML dataset -- see the Fairness section below (inherited unchanged from every contributing BP).

Recommendations

Re-run the platform rollup every time a BP's own report package regenerates

- **Specific:** This platform-wide rollup is a pure read-only pass-through of BP1-BP5's own real already-computed figures (see platform_rollup) -- it carries no independent validation of its own beyond the two structural/reconciliation gates (see this BP's own Notebook 3 report).
- **Measurable:** Re-run the extractor script + BP6 Notebooks 2-4 whenever any of BP1-BP5's MODEL_CARD.md/RULE_CARD.md or NB3 validation report changes; confirm Gate 2's real reconciliation check still passes exactly.
- **Achievable:** Uses the existing, already-built extractor script and BP6 pipeline -- no new infrastructure required.
- **Relevant:** Keeps the platform-wide figures this report cites from silently going stale relative to the per-BP source of truth.
- **Time-Bound:** Immediately after any BP1-BP5 report regeneration; otherwise on the same quarterly cadence as each BP's own recalibration review.

Reallocate the real platform-wide false-positive-reduction capacity gain

- **Specific:** Summed across BP1+BP4+BP5 (ASSUMPTION-labeled, never blended with category 2 or category 3 below): \$13,877,192 (\$13.88M) in false-positive-reduction savings, 213,496 investigator hours, 853,981 fewer false-positive alerts across the three transaction-level BPs.
- **Measurable:** Track real combined alert-review hours logged across BP1/BP4/BP5's production queues against this ASSUMPTION estimate.
- **Achievable:** Capacity freed is redeployed to case investigation depth across the three detection programs, not a headcount reduction -- an operational scheduling change.

- **Relevant:** Investigator capacity is the real, most commonly cited AML program bottleneck platform-wide (FFIEC BSA/AML Examination Manual, alert-management pillar).
- **Time-Bound:** Reassess real combined alert-review hours 90 days after platform-wide go-live.

Validate the real platform-wide true-positive uplift, led by BP4

- **Specific:** Summed across BP1+BP4+BP5: \$97,950,000 (\$97.95M) illustrative regulatory-exposure-avoidance, +1,959 additional real cases caught. BP4 contributes the single largest real category-2 dollar figure of the three (\$65,000,000 (\$65.00M)).
- **Measurable:** Real SAR filing rate on model-flagged alerts across the three contributing BPs, tracked per-BP and never blended with the false-positive savings in category 1 above.
- **Achievable:** Requires only tagging each filed SAR with which BP's model originally surfaced the alert -- a labeling change, not a new detection system.
- **Relevant:** Directs platform-wide model-risk review attention to BP4, the real largest real contributor to this benefit category.
- **Time-Bound:** First real platform-wide comparison report at the 6-month production mark.

Keep BP2's two typology-detection figures reported separately, never folded into categories 1/2

- **Specific:** BP2 contributes its own real auto-typing efficiency savings (\$1,313) and typology-confirmation value (\$1,104,000 (\$1.10M)) -- both real figures on a different unit basis (auto-typing hours, distinct case subset) than BP1/BP4/BP5's categories 1/2, per the locked never-blend-categories rule.
- **Measurable:** Any future dashboard or report citing a platform 'total benefit' figure is checked to confirm it reports these three categories separately, not as one blended number.
- **Achievable:** A documentation/labeling convention enforced by this report_builder.py module's own structural gate (see this BP's Notebook 3), not a manual check.
- **Relevant:** Prevents a materially misleading single 'platform savings' headline that would mix incompatible unit bases.
- **Time-Bound:** Immediate -- applies to this report and every future one referencing the platform rollout.

Document the real recurring champion algorithm (XGBoost) across the platform

- **Specific:** XGBoost is the real champion model on 3 of the platform's BPs (BP1, BP4, BP5), each independently selected by its own Stage A/B benchmark -- not assumed or forced to agree.
- **Measurable:** Confirm this cross-BP champion agreement holds at every future retrain; log explicitly if a future BP's own champion diverges.

- **Achievable:** Already computed by each BP's own existing Notebook 3 -- no new tooling required at the rollout level.
- **Relevant:** A champion algorithm stable across multiple independently-trained business problems is useful real evidence for platform-wide SR 11-7 model-risk documentation (shared infra/tooling/monitoring approach).
- **Time-Bound:** Refresh this cross-BP comparison at every platform-wide rollout regeneration.

Schedule a threshold review for BP1's real operating point (0.9798)

- **Specific:** BP1's real champion (XGBoost) operates at a real decision threshold of 0.9798 on LI-Medium, using 19 real engineered features, yielding real precision 0.142 and real recall 0.179 (real PR-AUC 0.1242).
- **Measurable:** Track real precision/recall drift against these two exact baseline figures on every production scoring batch; a real move of more than 5 percentage points on either triggers a threshold review.
- **Achievable:** Uses the already-built, self-tested BP1 scoring service -- no new infrastructure.
- **Relevant:** BP1 is the platform's foundational detection layer -- its threshold directly sets the real false-positive alert volume every downstream investigator queue inherits.
- **Time-Bound:** Quarterly, next due within 90 days of platform-wide go-live.

Grow real training examples for BP2's thinnest real typology class before the next retrain

- **Specific:** BP2's real class-count distribution on LI-Medium is uneven: 'GATHER-SCATTER' has 541 real matched examples versus 'RANDOM' at only 278 -- a real 1.9x gap, with real macro-F1 at 0.444 across 13 features.
- **Measurable:** Real per-class F1 (not just macro-F1) tracked every retrain; flag any real class whose F1 drops below the current run's weakest class.
- **Achievable:** Uses the existing LI-Medium/HI-Medium staged-scale-up pattern already locked in the master plan -- re-point DATASET_VARIANT, no new modeling approach.
- **Relevant:** A thin real typology class is the most likely place a future retrain silently regresses without a per-class check.
- **Time-Bound:** Before BP2's next scheduled retrain.

Re-confirm BP3's real bootstrap CI lower bound stays above 1.0x at the next scale-up

- **Specific:** BP3's real champion signal ('2-hop proximity to a TRAIN-flagged account') holds a real network-lift ratio of 3.19x on LI-Medium (2,032,095 real nodes, 4,363,197 real edges) -- the platform's only BP with no dollar figure (no sourced per-account cost basis, honestly disclosed rather than invented).
- **Measurable:** Real bootstrap 95% CI lower bound re-checked at every future real re-run; a lower bound at or below 1.0x (no real lift) is the pre-registered trigger for re-screening the candidate signal set.

- **Achievable:** Uses the existing, already-built Stage A/B screening and bootstrap CI pipeline -- no new infrastructure.
- **Relevant:** BP3's real lift is the only evidence its structural rule still outperforms random account selection -- losing it would mean investigators get no real benefit from network-aware triage.
- **Time-Bound:** At every future HI-Large/LI-Large stretch-validation attempt (master plan's own optional next step for this BP).

Investigate BP4's real precision/recall trade-off at threshold 0.9793

- **Specific:** BP4's real champion (XGBoost) achieves the platform's highest real recall among the three binary BPs (0.305) but its lowest real precision (0.086) at threshold 0.9793 on LI-Medium, using 24 real features (BP1's 19 plus BP4's own 5 structuring-specific ones).
- **Measurable:** Real precision/recall at a shifted threshold, re-evaluated against the SAME real held-out test set, before any production threshold change is made.
- **Achievable:** A real threshold-sweep re-evaluation using Notebook 3's already-saved test-set scores -- no retraining required.
- **Relevant:** BP4 contributes the platform's single largest real BENEFIT-category-2 dollar figure (\$65,000,000 (\$65.00M)) -- its real operating point deserves the platform's closest scrutiny.
- **Time-Bound:** Before BP4's real threshold is next changed in production.

Audit BP5's 27 real cross-border features for redundancy

- **Specific:** BP5's real champion (XGBoost) uses 27 real features -- the most of any binary BP on the platform -- at a real threshold of 0.9814, achieving real precision 0.176 and real recall 0.183 on LI-Medium.
- **Measurable:** Real SHAP global-importance ranking (already computed by BP5's own Notebook 3) reviewed for any cross-border feature contributing materially less than BP1's shared base features.
- **Achievable:** Uses BP5's already-computed SHAP output -- a review, not a new computation.
- **Relevant:** BP5 delivers the platform's single largest real BENEFIT-category-1 dollar figure (\$4,852,169 (\$4.85M)) -- confirming its extra features earn their real complexity protects that result.
- **Time-Bound:** Before BP5's feature set is next changed.

Keep BP3's real network-coverage volume as a portfolio-scale line, never dollarized or summed with BENEFIT

- **Specific:** BP3's real network covers 2,032,095 nodes and 4,363,197 edges -- structurally distinct from BP1/BP2/BP4/BP5's transaction-level populations, with no sourced per-account investigation-cost basis to dollarize it (disclosed honestly, never invented).
- **Measurable:** Any future platform dashboard reporting a single 'accounts covered' or 'network scale' figure cites this real count directly, never converts it to a dollar estimate.

- **Achievable:** A reporting convention, not a new computation -- BP3's own RULE_CARD.md already discloses the same real no-cost-basis limitation.
- **Relevant:** Prevents a fabricated per-account cost figure from entering platform-wide financial reporting.
- **Time-Bound:** Immediate -- applies to this report and every future one referencing BP3's portfolio scale.

Maintain the real five-BP all-PASS standard on every future platform-wide retrain cycle

- **Specific:** All five real business problems currently pass both validation gates (BP1, BP2, BP3, BP4, BP5).
- **Measurable:** Every BP must continue to PASS both gates before its figures feed the next platform rollup regeneration.
- **Achievable:** Enforced automatically by each BP's own existing Notebook 3 gate logic plus BP6 Notebook 3's own reconciliation gate.
- **Relevant:** This is the basis for this report's current platform-wide recommended status.
- **Time-Bound:** Every retrain cycle, before the next rollup regeneration.